



**Department of Electrical and Computer Engineering
North South University**

Senior Design Project
**Multi-Perspective Reasoning Evaluation of QLoRA-
Fine-Tuned Small Language Models**

Md. Miraz Ihsan Chowdhury ID# 2131407642

Abdus Samad Sunny ID# 2212047042

Faculty Advisor:

Dr. Mohammad Ashrafuzzaman Khan

Associate Professor

ECE Department

Spring, 2026

LETTER OF TRANSMITTAL

May 5, 2026

To

Dr. Mohammad Abdul Matin
Chairman,
Department of Electrical and Computer Engineering,
North South University, Dhaka.

Subject: Submission of Capstone Project Report on “Multi-Perspective Reasoning Evaluation of QLoRA-Fine-Tuned Small Language Models”

Dear Sir,

With due respect, we would like to submit our **Capstone Project Report** on “**Multi-Perspective Reasoning Evaluation of QLoRA-Fine-Tuned Small Language Models**” as a part of our BSc program. This project involves fine-tuning four small open-source Large Language Models — Phi-2, Phi-3 Mini, Qwen 2.5 Math 7B, and DeepSeekMath 7B — on the Thinkless dataset using QLoRA. The goal was to train these models to generate structured reasoning-based answers. After fine-tuning, we ran inference using a `generate_answer()` function and observed how well each model produced structured outputs in the format of Initial Solution, Reflections, and a final Synthesized answer. This project was a highly valuable learning experience, providing hands-on exposure to fine-tuning small language models, evaluating their reasoning outputs, and analyzing their failure modes. We have tried to the best of our ability to meet all the academic and technical dimensions required by this report.

We will be highly obliged if you kindly receive this report and provide your valuable judgment. It would be our immense pleasure if you find this report useful and informative to have an apparent perspective on the issue.

Sincerely Yours,

.....
Md. Miraz Ihsan Chowdhury
ECE Department
North South University, Bangladesh

.....
Abdus Samad Sunny
ECE Department
North South University, Bangladesh

APPROVAL

Md. Miraz Ihsan Chowdhury (ID # 2131407642) and Abdus Samad Sunny (ID #2212047042) from Electrical and Computer Engineering Department of North South University, have worked on the Senior Design Project titled “**Multi-Perspective Reasoning Evaluation of QLoRA-Fine-Tuned Small Language Models**” under the supervision of Dr. Mohammad Ashrafuzzaman Khan partial fulfillment of the requirement for the degree of Bachelors of Science in Engineering and has been accepted as satisfactory.

Supervisor’s Signature

.....

Dr. Mohammad Ashrafuzzaman Khan

Associate Professor

Department of Electrical and Computer Engineering
North South University
Dhaka, Bangladesh.

Chairman’s Signature

.....

Dr. Mohammad Abdul Matin

Professor

Department of Electrical and Computer Engineering
North South University
Dhaka, Bangladesh.

DECLARATION

This is to declare that this project is our original work. No part of this work has been submitted elsewhere partially or fully for the award of any other degree or diploma. All project related information will remain confidential and shall not be disclosed without the formal consent of the project supervisor. Relevant previous works presented in this report have been properly acknowledged and cited. The plagiarism policy, as stated by the supervisor, has been maintained.

Students' names & Signatures

1. Md. Miraz Ihsan Chowdhury

2. Abdus Samad Sunny

ACKNOWLEDGEMENTS

The authors would like to express their heartfelt gratitude towards their project and research supervisor, Dr. Mohammad Ashrafuzzaman Khan, Assistant Professor, Department of Electrical and Computer Engineering, North South University, Bangladesh, for his invaluable support, precise guidance and advice pertaining to the experiments, research and theoretical studies carried out during the course of the current project and also in the preparation of the current report.

Furthermore, the authors would like to thank the Department of Electrical and Computer Engineering, North South University, Bangladesh for facilitating the research. We would also like to thank Ahmed Faizul Haque Dhrubo for helping us in this project. The authors would also like to thank their loved ones for their countless sacrifices and continual support.

ABSTRACT

Multi-Perspective Reasoning Evaluation of QLoRA-Fine-Tuned Small Language Models

Large Language Models (LLMs) have demonstrated strong performance on mathematical reasoning tasks, yet deploying them on consumer-grade hardware remains a significant challenge due to their high computational demands. The Thinkless framework (2025) introduced a hybrid reasoning dataset in which model responses follow a structured multi-stage format — comprising an initial solution, multiple reflective perspectives, and a final synthesis. In this work, we fine-tune four small open-source LLMs — Phi-2 (2.7B), Phi-3 Mini (3.8B), Qwen 2.5 Math 7B, and DeepSeek-Math 7B — on this dataset using QLoRA, within Kaggle's free-tier GPU environment. After fine-tuning, we evaluate each model's ability to produce coherent, structured multi-perspective reasoning outputs through a dedicated inference pipeline. Our results reveal a notable divergence across models: Qwen 2.5 Math 7B and Phi-2 consistently produced correct and well-structured answers, whereas Phi-3 Mini exhibited repetitive and unstable generation patterns, and DeepSeek-Math 7B failed to converge due to tokenizer incompatibility with the training format. These findings suggest that a model's pre-training domain alignment is a stronger predictor of fine-tuning success than parameter count alone, offering practical insights into the challenges and trade-offs of adapting small LLMs for structured mathematical reasoning.

TABLE OF CONTENTS

LETTER OF TRANSMITTAL	2
APPROVAL	4
DECLARATION	5
ACKNOWLEDGEMENTS.....	6
ABSTRACT.....	7
LIST OF FIGURES	10
LIST OF TABLES.....	11
Chapter 1 Introduction	12
1.1 Background and Motivation	12
1.2 Purpose and Goal of the Project.....	12
1.3 Organization of the Report.....	14
Chapter 2 Research Literature Review	15
2.1 Existing Research and Limitations.....	15
Chapter 3 Methodology	16
3.1 System Design.....	16
3.2 Hardware and/or Software Components	17
3.3 Hardware and/or Software Implementation	17
Chapter 4 Investigation/Experiment, Result, Analysis and Discussion.....	19
Chapter 5 Impacts of the Project.....	22
5.1 Impact of this project on societal, health, safety, legal and cultural issues.....	22
5.2 Impact of this project on environment and sustainability	22
Chapter 6 Project Planning and Budget	24
Chapter 7 Complex Engineering Problems and Activities	25
7.1 Complex Engineering Problems (CEP).....	25
7.2 Complex Engineering Activities (CEA)	26
Chapter 8 Conclusions	27

8.1 Summary	27
8.2 Limitations	27
8.3 Future Improvement.....	27
References.....	28

LIST OF FIGURES

Figure 1. A sample Gantt chart.
not defined.

Error! Bookmark not defined.

Error! Bookmark

LIST OF TABLES

TABLE I. A SAMPLE SOFTWARE/HARDWARE TOOLS TABLE	14
TABLE II. A SAMPLE COMPLEX ENGINEERING PROBLEM ATTRIBUTES TABLE	18
TABLE III. A SAMPLE COMPLEX ENGINEERING PROBLEM ACTIVITIES TABLE	19

Chapter 1 Introduction

1.1 Background and Motivation

The surge in Large Language Models (LLMs) has revolutionized natural language processing, but the internal mechanisms of their "reasoning" processes remain largely opaque and computationally expensive. Recent breakthroughs in Large Reasoning Models (LRMs), such as OpenAI's o1 and DeepSeek-R1 [6], have demonstrated that scaling inference-time compute through Chain-of-Thought (CoT) prompting can significantly enhance logical performance [2]. However, these models typically rely on massive parameter counts (70B+), making them inaccessible for real-time edge deployment and consumer-grade hardware.

Small Language Models (SLMs), such as the Phi and Qwen families, offer a more sustainable and privacy-conscious path for local deployment. Despite their efficiency, the logical "thinking" depth of SLMs is often questioned, as they frequently struggle with the structural stability required for multi-stage deliberation. This project is motivated by the critical need to bridge this gap: to understand if models with sub-10B parameters can emulate the complex, multi-perspective reasoning found in their larger counterparts without suffering from catastrophic forgetting or repetitive instability. By leveraging Parameter-Efficient Fine-Tuning (PEFT) techniques like QLoRA [5] and the adaptive reasoning framework introduced in the Thinkless (2025) study [1], this research explores whether architectural alignment and high-quality "thinking" data can compensate for reduced parameter scale.

1.2 Purpose and Goal of the Project

The core objective of this project is to investigate the feasibility of inducing complex, multi-stage reasoning capabilities in Small Language Models (SLMs) through structured fine-tuning. While large-scale models have demonstrated robust logical reasoning, the ability of sub-10B parameter models to sustain a coherent "Chain-of-Thought" across multiple perspectives remains under-explored.

Primary Objectives

- **Feasibility Assessment of QLoRA:** To determine if 4-bit quantized Low-Rank Adaptation (QLoRA) can effectively train small models (2.7B to 7B) to adhere to strict reasoning formats without losing general linguistic capabilities.

- **Multi-Stage Deliberation:** To develop and implement an inference pipeline that mirrors human-like deliberation, involving an initial solution, critical reflection, and synthesis.
- **Architectural Benchmarking:** To conduct a rigorous comparative study across four distinct architectures (Phi-2, Phi-3, Qwen 2.5, and DeepSeek) to identify which design traits favor logical stability.

Contributions of the Project

The major contributions of this work are as follows:

1. **Reasoning Adapter Suite:** We provide a set of four fine-tuned LoRA adapters optimized for the Thinkless hybrid-reasoning format, enabling small models to utilize <think> tags effectively.
2. **Diagnostic Failure Matrix:** We contribute a detailed analysis of "Inference Instability" in small models, specifically documenting the repetition loops in Phi-3 Mini and the tokenizer-level failures in DeepSeek-Math.
3. **Domain Alignment Insights:** We provide empirical evidence that a model's pre-training domain (e.g., Qwen's mathematical focus) is a more critical success factor for reasoning fine-tuning than the absolute parameter count.

Novelty of the Work

The novelty of this project lies in its "Multi-Perspective Deliberation" probe. While most existing research focuses on "Single-Shot" Chain-of-Thought, our work is among the first to force sub-5B models (like Phi-2) to engage in self-reflection and multi-path synthesis. By attempting to port the high-level logic of the Thinkless (2025) framework into miniature architectures, we challenge the current assumption that structured, adaptive thinking is a property exclusive to Large Language Models. This study provides a unique "stress test" for small models, exposing the boundary where mimicry of a reasoning format ends and genuine logical progression begins.

1.3 Organization of the Report

The remainder of this report is organized into eight chapters, documenting the progression from theoretical research to technical implementation and final analysis:

- Chapter 2 (Research Literature Review): Evaluates existing literature on Chain-of-Thought (CoT) reasoning, parameter-efficient fine-tuning (PEFT), and the current state of small language model architectures.
- Chapter 3 (Methodology): Provides a detailed description of the system architecture, including the software stack, the QLoRA training pipeline, and the iterative multi-perspective inference engine.
- Chapter 4 (Investigation, Result, Analysis and Discussion): Discusses the experimental setup, presents the training metrics and inference results, and provides an in-depth analysis of the performance differences between the tested models.
- Chapter 5 (Impacts of the Project): Analyzes the project's multi-dimensional influence on societal, health, safety, legal, and environmental sustainability issues.
- Chapter 6 (Project Planning and Budget): Outlines the project timeline through a Gantt chart and provides an economic analysis of the resources utilized.
- Chapter 7 (Complex Engineering Problems and Activities): Documents the specific engineering challenges (CEP) and activities (CEA) encountered during the model adaptation and evaluation phases.
- Chapter 8 (Conclusions): Summarizes the findings of the study, identifies project limitations, and proposes directions for future research and technical improvements.
- References: Lists the academic sources and technical reports cited throughout the project using the IEEE format.

Chapter 2 Research Literature Review

2.1 Existing Research and Limitations

Recent research has focused extensively on the development of reasoning-capable language models and efficient training methodologies. VainF et al. [1] developed the Thinkless framework, an adaptive reasoning system that allows models to switch between <short> and <think> modes. The authors utilized a reinforcement learning paradigm to demonstrate that reasoning models can reduce computational cost by 50–90% while maintaining high accuracy on mathematical benchmarks. J. Wei et al. [2] established the concept of Chain-of-Thought (CoT), showing that eliciting step-by-step logic through intermediate steps significantly improves the logical performance of LLMs. To facilitate the training of such models on constrained hardware, T. Dettmers et al. [5] implemented QLoRA, a 4-bit quantization technique that enables high-fidelity fine-tuning with minimal VRAM requirements. Furthermore, Microsoft Research [12] presented Phi-2, a 2.7B parameter model that demonstrated that high-quality synthetic data could allow small models to outperform much larger architectures on reasoning and logic tasks.

The following observations have been made after a detailed examination of the literature reviews: (i) most of the articles focused on scaling reasoning capabilities in large-scale models (70B+), leaving the stability of sub-10B models under-explored, (ii) in general, the articles did not compare the performance of multiple distinct small architectures (such as Phi, Qwen, and DeepSeek) when subjected to the same structured reasoning training, and (iii) there is an absence of evaluating multi-stage deliberation (reflection and synthesis) in miniature models.

These investigations have motivated us to implement a Multi-Perspective Reasoning Evaluation system in this project using the Thinkless hybrid dataset. We apply the QLoRA fine-tuning approach across four distinct architectures (Phi-2, Phi-3 Mini, Qwen 2.5 Math, and DeepSeek-Math) to compare their logical coherence. Finally, we implement a multi-stage inquiry pipeline consisting of initial solutions, reflective perspectives, and synthesized answers to stress-test the reasoning depth of these small-scale language models.

Chapter 3 Methodology

3.1 System Design

The system architecture for this project is designed as an integrated framework that combines parameter-efficient fine-tuning with an iterative deliberation pipeline to enable complex reasoning on consumer-grade hardware. The design begins with a Quantized Model Core, where base architectures such as Phi, Qwen, and DeepSeek are compressed using 4-bit NormalFloat (NF4) quantization via the bitsandbytes library to ensure memory efficiency. Above this core, the system utilizes LoRA (Low-Rank Adaptation) adapters targeting the attention projection layers to learn the specific hybrid reasoning syntax of the Thinkless dataset. The operational logic is governed by an Iterative Multi-Perspective Inference Engine, which replaces standard single-pass generation with a structured refinement loop. In this workflow, a central controller manages a sequence of stages: an initial solution is generated, followed by an iterative loop of self-reflections that critique the prior logic, and finally, a synthesis stage that aggregates all perspectives into a verified final answer. This entire process is supported by a dynamic context buffer that sequentially appends internal monologues to the prompt, ensuring the final output is grounded in a holistic deliberation. The system flow follows a logical progression from raw data preprocessing and QLoRA training to the final multi-stage inference execution, providing a robust end-to-end pipeline for evaluating the reasoning depth of small language models.




Figure: System Design

3.2 Hardware and/or Software Components

The implementation of this project utilizes a combination of high-performance cloud computing and specialized deep learning libraries designed for parameter-efficient fine-tuning. The hardware setup is centered around Kaggle’s Tesla T4 GPU environment, which provides the necessary 16GB of VRAM to host 7B parameter models in 4-bit quantized mode. The software stack is built on the PyTorch ecosystem, utilizing the HuggingFace suite for model management, PEFT for LoRA implementation, and BitsAndBytes for 4-bit quantization.

TABLE I. A SAMPLE SOFTWARE/HARDWARE TOOLS TABLE

Tool	Functions	Other similar Tools (if any)	Why selected this tool
Kaggle Tesla T4	Cloud GPU Hardware	Google Colab, AWS	Provides free 16GB VRAM, essential for fine-tuning 7B models.
PyTorch	Deep Learning Framework	TensorFlow, JAX	Industry standard with extensive support for HuggingFace and PEFT.
HuggingFace Transformers	Model & Tokenizer Management	PyText, DeepSpeed	Offers the most comprehensive library of pre-trained SLMs (Phi, Qwen, etc.).
BitsAndBytes	4-bit Quantization (NF4)	AutoGPTQ, AWQ	Seamlessly integrates with Transformers for QLoRA implementation.
PEFT (LoRA)	Parameter-Efficient Tuning	Adapters, Prompt Tuning	Allows fine-tuning 7B models with only ~0.2% trainable parameters.
Datasets (HF)	Data Handling & Cleaning	Pandas, Dask	Optimized for large-scale NLP datasets and fast tokenization.
Python	Programming Language	R, C++	The primary language for AI research with the most robust library ecosystem.

3.3 Hardware and/or Software Implementation

The software implementation is structured as a continuous pipeline that begins with the quantized initialization of base models using BitsAndBytesConfig, enabling 4-bit

NormalFloat (NF4) precision to fit large architectures into limited GPU memory. Following this, LoRA (Low-Rank Adaptation) adapters are integrated into the model's attention layers using the PEFT library, which allows for the training of high-rank reasoning logic while keeping the billions of base parameters frozen. The data layer is implemented through a custom tokenization and formatting module that merges short-form and long-form responses from the Thinkless dataset into a unified training template, ensuring the model learns to associate specific queries with structured <think> blocks. The final phase involves the development of an Iterative Inference Engine via the `multi_reasoning_workflow` function; this module acts as a state machine that sequentially captures initial thoughts, generates critical reflections by injecting prior context back into the prompt, and executes a final synthesis stage to produce a verified answer. This integrated software stack ensures that the model's learned reasoning capabilities are effectively utilized through a deliberative, multi-stage generation process.

Chapter 4 Investigation/Experiment, Result, Analysis and Discussion

Experimental Parameters and Variables

The experiment was designed to measure the impact of multi-perspective fine-tuning on four distinct SLM architectures. The Independent Variable was the model architecture (Phi vs. Qwen vs. DeepSeek), and the Dependent Variables were the training loss convergence and the logical consistency of the multi-stage outputs.

TABLE. TRAINING CONFIGURATIONS ACROSS MODELS

Model	Parameters	Training Loss (Final)	Validation Loss (Final)	Status
Phi-2	2.7B	0.954	1.121	Successful
Phi-3 Mini	3.8B	0.072	1.782	Overfitted
Qwen 2.5 Math	7B	0.421	0.634	Successful
DeepSeek-Math	7B	N/A	N/A	Tokenizer Failure

Training Result Analysis

As shown in Table VI, the models exhibited divergent learning behaviors.

Phi-2 showed the most balanced learning curve, maintaining a narrow gap between training and validation loss, suggesting that the 2.7B architecture is highly receptive to LoRA-based reasoning adapters.

Qwen 2.5 Math achieved the lowest stable validation loss (0.634). Its pre-existing mathematical weights significantly accelerated the adaptation to the Thinkless reasoning format.

Phi-3 Mini displayed a "loss collapse" (0.072), which resulted in severe Repetition Loops during inference. This indicated that the model had memorized the training format without internalizing the logic.

Qualitative Result: The Multi-Perspective Evaluation

To test the "Inquiry Pipeline," we used the "Lois's Books" logic puzzle. The results highlight the effectiveness of the Reflection-Synthesis stages.

TABLE. INFERENCE PERFORMANCE ON LOGIC PUZZLES

Model	Initial Solution	Reflection 1	Synthesis	Result
Qwen 2.5 Math	Calculated $20 + x = 23$.	Attempted working backwards.	Corrected the count of donated books.	Correct
Phi-2	Identified sibling ratios correctly.	Verified the family tree.	Simplified the answer to 1 sister.	Correct
Phi-3 Mini	Started strong logic.	Repeated the Initial Solution.	Failed to produce a final answer.	Failed

Constructive Discussion

The investigation reveals that Model Size is not the primary driver of reasoning success in SLMs. Despite being the smallest model at 2.7B, Phi-2 outperformed the 3.8B Phi-3 Mini and matched the 7B models in logical consistency. This suggests that the Phi-2 architecture possesses a superior "inductive bias" for structured reasoning when paired with QLoRA.

The failure of DeepSeek-Math highlights a critical engineering problem: Special Token Incompatibility. Because DeepSeek utilizes a unique tokenizer structure, the special tags used in the Thinkless dataset triggered an EOS (End of Sentence) token prematurely, preventing the multi-stage generation.

Conclusion of the Experiment

The results confirm that a Multi-Perspective Inquiry Pipeline significantly improves the performance of fine-tuned SLMs. Qwen 2.5 Math demonstrated that a model can "self-correct" its initial hallucinations during the reflection phase, a behavior typically seen only in models like DeepSeek-R1 or GPT-4. This proves that high-quality reasoning is achievable on consumer-grade hardware provided the model's pre-training domain is well-aligned with the task.

Chapter 5 Impacts of the Project

5.1 Impact of this project on societal, health, safety, legal and cultural issues

The development of efficient reasoning capabilities in small-scale language models has a multi-dimensional impact across various sectors:

- **Societal and Cultural:** This project promotes the democratization of artificial intelligence by enabling high-level reasoning on low-cost, consumer-grade hardware. Culturally, this shifts the power dynamic away from centralized "Big Tech" corporations, allowing local communities and developers in diverse regions to build and deploy AI that reflects their own values and languages without the high entry barrier of massive computational costs.
- **Health:** By facilitating on-device reasoning, this project has a profound impact on the healthcare sector. Patients can use AI-driven diagnostic tools or mental health assistants where their sensitive medical data never leaves their local device, ensuring absolute confidentiality and fostering trust in digital health interventions.
- **Safety:** The multi-perspective reasoning framework implemented in this project directly improves the reliability and safety of AI systems. By forcing the model to reflect and synthesize multiple views, the system significantly reduces the "hallucination" rate, making it safer for high-stakes applications like academic tutoring, financial calculations, and structural design.
- **Legal:** From a legal perspective, the local-first approach used in this project simplifies compliance with global data privacy laws, such as the General Data Protection Regulation (GDPR). Since the reasoning process is decentralized, it removes the legal liabilities and risks associated with data breaches and unauthorized third-party access during cloud transmission.

5.2 Impact of this project on environment and sustainability

The fine-tuning of small language models (SLMs) using parameter-efficient techniques like QLoRA carries significant implications for environmental impact as it contributes to the

massive reduction of the carbon footprint associated with AI development. While training large-scale models (70B+ parameters) consumes vast amounts of electricity and generates significant heat, this project demonstrates that high-level reasoning can be achieved using exponentially less power on consumer-grade GPUs. On the other hand, this work holds notable ecological and sustainability significance as it promotes the use of existing, lower-end hardware, thereby extending the lifecycle of electronic devices and contributing to a reduction in electronic waste (e-waste). These projects demonstrate the potential of efficient AI technology to mitigate the environmental challenges of the "big data" era and promote a more sustainable, energy-conscious future for artificial intelligence.

Chapter 6 Project Planning and Budget

The project is structured across two terms, spanning from January 22, 2026, to August 20, 2026. As illustrated in Figure 1, the first term focuses on Research and Development, encompassing the literature review, dataset preparation, and the initial QLoRA training of the four selected architectures (Phi-2, Phi-3, Qwen 2.5, and DeepSeek). The second term is dedicated to Evaluation and Analysis, where the models undergo rigorous performance benchmarking, error analysis, and efficiency testing. This systematic approach ensures that each technical milestone—from prompt engineering to system integration—is achieved within the designated timeframe.




Figure: Gantt chart.

Chapter 7 Complex Engineering Problems and Activities

7.1 Complex Engineering Problems (CEP)

The implementation of multi-perspective reasoning in resource-constrained models involves several complex engineering challenges that span algorithmic design, hardware optimization, and data science. The following table summarizes how these attributes were addressed:

TABLE II. A SAMPLE COMPLEX ENGINEERING PROBLEM ATTRIBUTES TABLE

Attributes		Addressing the complex engineering problems (P) in the project
P1	Depth of knowledge required (K3-K8)	The project requires deep knowledge of Neural Networks (K3), Natural Language Processing (K4), Fine-tuning/Optimization Algorithms (K5), PEFT & Quantization Tools (K6), Computational Energy Efficiency (K7), and Contemporary Research in LLM Reasoning (K8).
P2	Range of conflicting requirements	A significant conflict exists between model size (which dictates reasoning depth) and VRAM capacity (which limits model loading). Additionally, there is a trade-off between inference latency and the number of reflection cycles required for high accuracy.
P3	Depth of analysis required	There is no unique path to inducing logic in SLMs. Depth of analysis was required to select the optimal fine-tuning strategy (QLoRA vs. LoRA) and to tune hyperparameters across four distinct and competing architectures.
P4	Familiarity of issues	Involves complex configurations of Transformer architectures, 4-bit NormalFloat (NF4) quantization, LoRA adapters, and cloud-based GPU orchestration.
P5	Extent of applicable codes	While no formal governmental code exists, the project strictly adheres to Open-Source Licenses (Apache 2.0/MIT) and Ethical AI Guidelines regarding the mitigation of model hallucinations.
P6	Extent of stakeholder involvement	Stakeholders include the global NLP research community, developers of edge-based AI applications, academic educators, and students who utilize SLMs for local tutoring.
P7	Interdependence	The project involves highly interdependent sub-systems, including tokenizer mapping, quantized weight matrices, LoRA adapter integration, and the iterative deliberation logic of the inference engine.

7.2 Complex Engineering Activities (CEA)

The execution of this project required a wide range of engineering activities, from managing high-performance cloud resources to implementing innovative reasoning algorithms. These activities are summarized in the table below:

TABLE III. A SAMPLE COMPLEX ENGINEERING PROBLEM ACTIVITIES TABLE

Attributes		Addressing the complex engineering activities (A) in the project
A1	Range of resources	The project involves cloud-based GPU hardware (Tesla T4), modern AI software stacks (PyTorch, PEFT), high-dimensional reasoning datasets (Thinkless), and significant human research and programming effort.
A2	Level of interactions	Involves technical interactions between group members to synchronize model fine-tuning, collaboration with the research supervisor for methodological refinement, and engagement with the open-source AI community for troubleshooting.
A3	Innovation	Employs innovative engineering skills by porting the complex adaptive thinking framework into miniature architectures and developing a multi-perspective synthesis pipeline to benchmark the limits of small-scale logic.
A4	Consequences to society / Environment	Significant impact on environmental sustainability by reducing the energy required for AI training. Contributes to social equity by providing high-quality logical reasoning tools that can run on low-cost hardware.
A5	Familiarity	Requires advanced familiarity with Transformer architectures, QLoRA quantization, and recursive prompt engineering. Aligns with UN SDG #4 (Quality Education) through accessible AI tutors and UN SDG #9 (Industry, Innovation, and Infrastructure).

Chapter 8 Conclusions

8.1 Summary

Our study confirms that small models can indeed perform structured reasoning, but success depends heavily on the pre-training data. Models pre-trained specifically on math/code (Qwen Math) performed significantly better than general-purpose models (Phi-3).

8.2 Limitations

The primary limitation was the hardware constraint of free GPU environments, which forced a reduction in training sample size (0.2% of the original dataset).

8.3 Future Improvement

Future work will focus on three key areas to enhance the reasoning depth of small models:

- **Preference Optimization:** Implementing DPO or RLHF to penalize the repetition loops observed in Phi-3 and reward logical coherence.
- **Architectural Scalability:** Resolving the tokenizer mapping issues for DeepSeek models and scaling the training from a subset to the full 1M+ Thinkless dataset.
- **Adaptive Deliberation:** Developing a dynamic reflection mechanism where models decide the number of thinking cycles based on problem difficulty.
- **Quantitative Benchmarking:** Conducting large-scale testing on GSM8K and MATH datasets to provide rigorous statistical validation of the fine-tuned adapters.

References

- [1] VainF, et al., "Thinkless: LLM Learns When to Think," *arXiv preprint arXiv:2505.13379*, 2025.
- [2] J. Wei, et al., "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [3] X. Wang, et al., "Self-Consistency Improves Chain of Thought Reasoning in Language Models," in *International Conference on Learning Representations (ICLR)*, 2023.
- [4] E. J. Hu, et al., "LoRA: Low-Rank Adaptation of Large Language Models," in *International Conference on Learning Representations (ICLR)*, 2022.
- [5] T. Dettmers, A. Pagnoni, A. Holtzman, and L. Zettlemoyer, "QLoRA: Efficient Finetuning of Quantized LLMs," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [6] DeepSeek-AI, "DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning," *arXiv preprint arXiv:2501.12948*, 2025.
- [7] M. Abdin, et al., "Phi-3 Technical Report: A Highly Capable Language Model Locally on Your Phone," *arXiv preprint arXiv:2404.14219*, 2024.
- [8] Qwen Team, "Qwen2.5-Math Technical Report: Toward Mathematical Expert at One-Billion Scale," *arXiv preprint arXiv:2409.12122*, 2024.
- [9] K. Cobbe, et al., "Training Verifiers to Solve Math Word Problems," *arXiv preprint arXiv:2110.14168*, 2021.
- [10] H. Touvron, et al., "Llama 2: Open Foundation and Fine-Tuned Chat Models," *arXiv preprint arXiv:2307.09288*, 2023.
- [11] A. Q. Jiang, et al., "Mistral 7B," *arXiv preprint arXiv:2310.06825*, 2023.
- [12] Microsoft Research, "Textbooks Are All You Need II: Phi-2 Technical Report," *Microsoft Blog*, 2023.